



Amazon Rainforest Image Classification



CONNOR MCANUFF - NOVEMBER 2019

Content

Slide Index

Project Overview:	3-5
Data Wrangling:	6-8
Exploratory Data Analysis:	9-24
Machine Learning:	25-45
Conclusions:	46-48

Project Overview – Amazon Rainforest

- **6,000,000 km²** across multiple countries.
- Several million species of plant, tree, insect, and animal life.
- **Contains human civilization** resulting in encroachment, exploitation, deforestation, and other forms of destruction:
 - **Reduced biodiversity;**
 - **Habitat loss;**
 - **Climate change;**
 - **Desertification;**
 - **Soil erosion.**

Project Overview – Satellite Imagery

- Can be used to monitor large areas of land.
- Previously, 30 m/pixel or 250 m/pixel imagery used – cannot identify small-scale deforestation/degradation.
- **Planet (company) plans to collect daily 3-5 m/pixel imagery of Earth's surface.**



Project Overview – Value in Classifying Satellite Images

- Scale and frequency of the Amazon Rainforest imaging means it would be **extremely cumbersome and expensive to manually classify images**.
- **Automated classification** would allow for **rapid mapping** of Amazon Rainforest to identify deforestation/degradation – action can be taken by governments or NGOs.
- Other uses such as infrastructure, aid, census, and resource planning.

Data Wrangling – Dataset Overview

- Sourced from former Kaggle competition.
- **40,479 satellite images:**
 - 256 x 256 pixels (946.2m x 947.2m on ground).
 - TIFF (4-channel) and JPEG (3-channel) format.
- **17 classifications (multilabels):**
 - Atmospheric conditions:
 - *Cloudy, partly cloudy, hazy.*
 - Common land cover/use:
 - *Primary rainforest, water, habitation, agriculture, road, cultivation, bare ground.*
 - Uncommon land cover/use:
 - *Slash and burn, conventional mining, selective logging, artisanal mining, blooming, blow down.*

Data Wrangling – Dataset Overview (2)


- CSV file with list of image file names and their associated labels/tags:

image_name	tags
train_0	haze primary
train_1	agriculture clear primary water
train_2	clear primary
train_3	clear primary
train_4	agriculture clear habitation primary road
train_5	haze primary water
train_6	agriculture clear cultivation primary water


Data Wrangling – Importing and Cleaning

- CSV file read directly into a Pandas DataFrame.
- Images brought in for viewing using opencv methods:
 - imread
 - imshow
- Dataset file constructed for use in machine learning models.


Exploratory Data Analysis – Tag Occurrences




Exploratory Data Analysis – Tag Co-occurrences



Exploratory Data Analysis – Image Samples




train_10 JPEG



train_10 TIFF



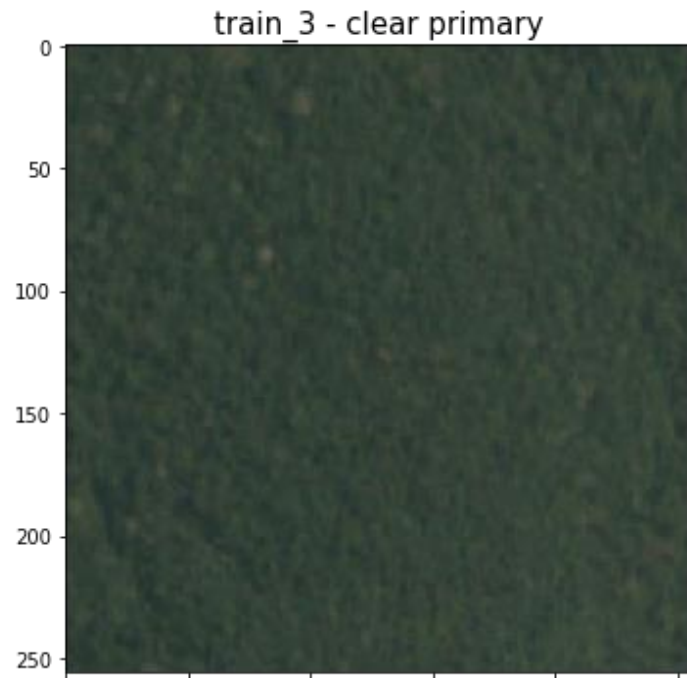
Exploratory Data Analysis – TIFF Images



Exploratory Data Analysis – Image Samples

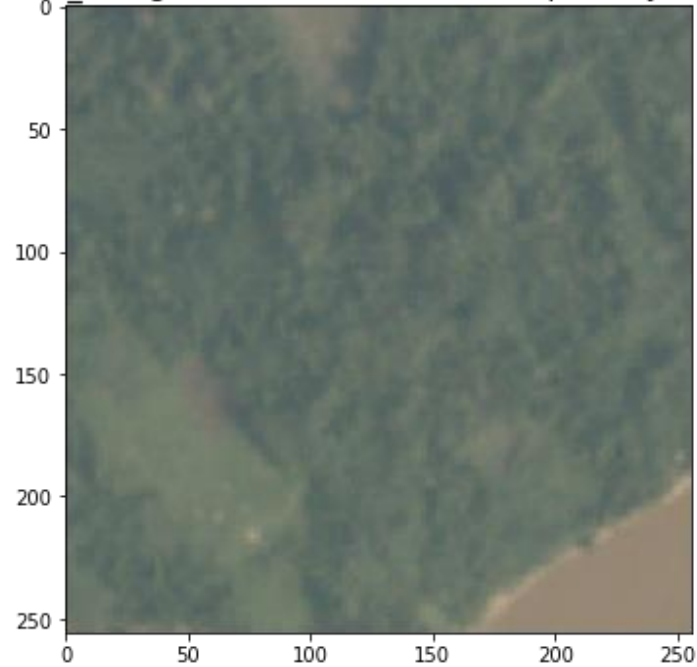


Exploratory Data Analysis – Image Samples (2)




Exploratory Data Analysis – Image Samples (3)

train_6 - agriculture clear cultivation primary water



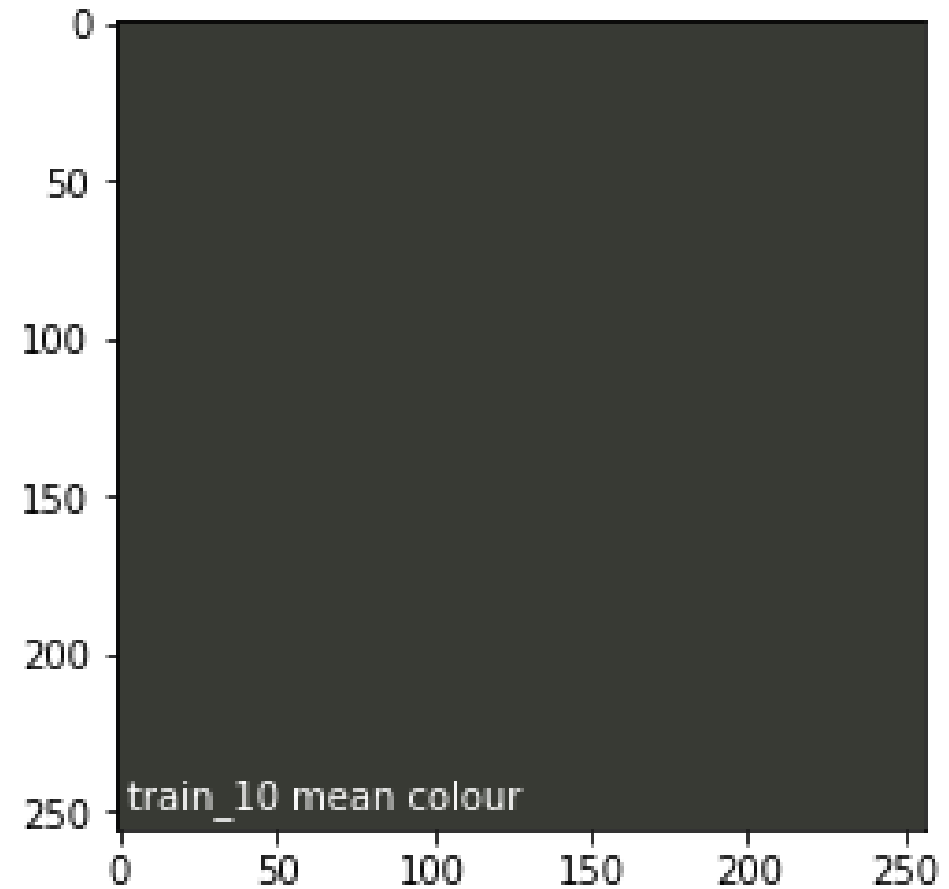

train_7 - haze primary




train_8 - agriculture clear cultivation primary



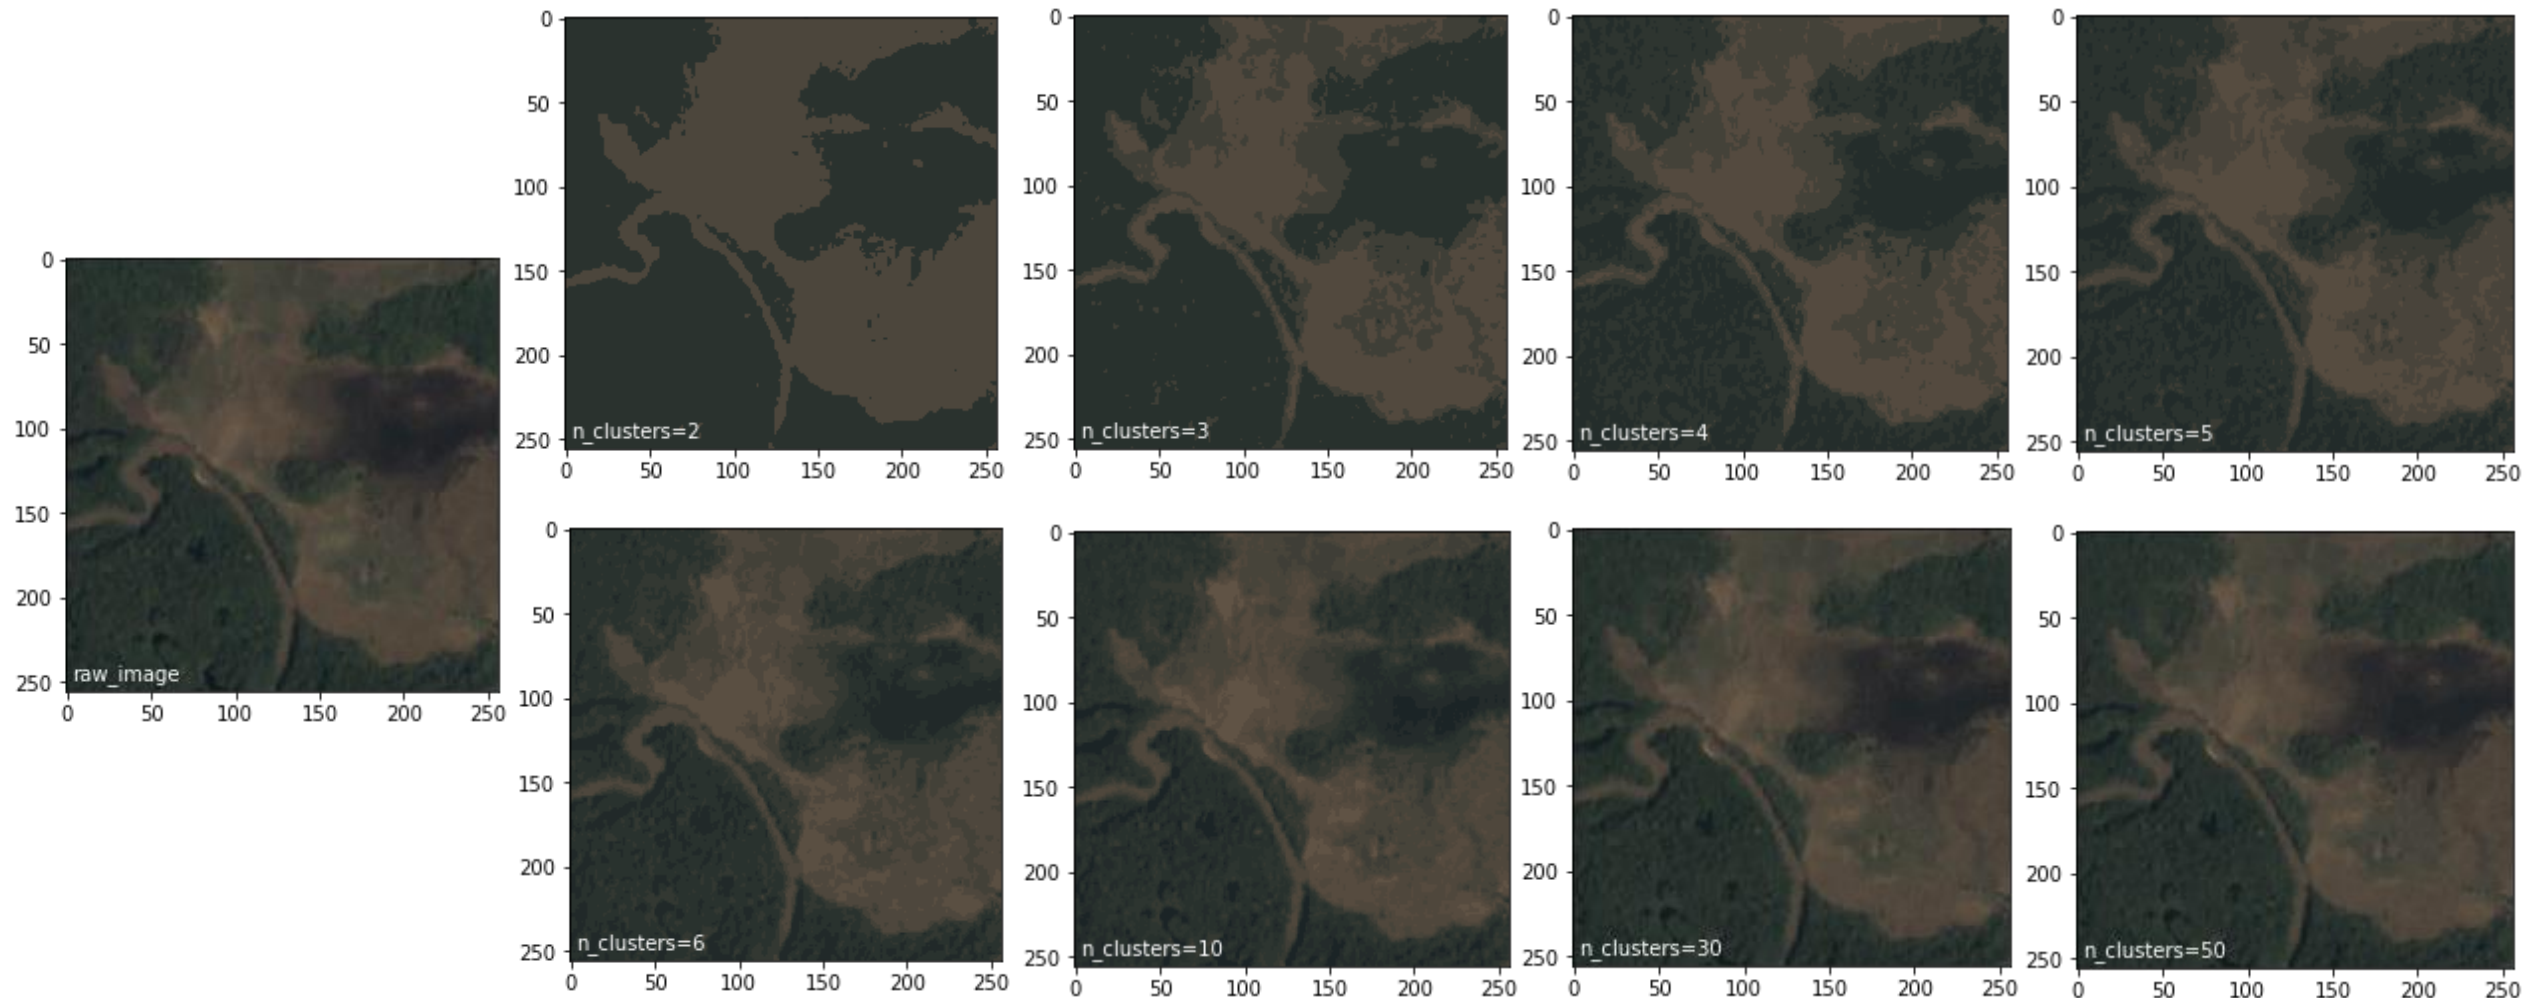
Exploratory Data Analysis – Single Image Mean Colour



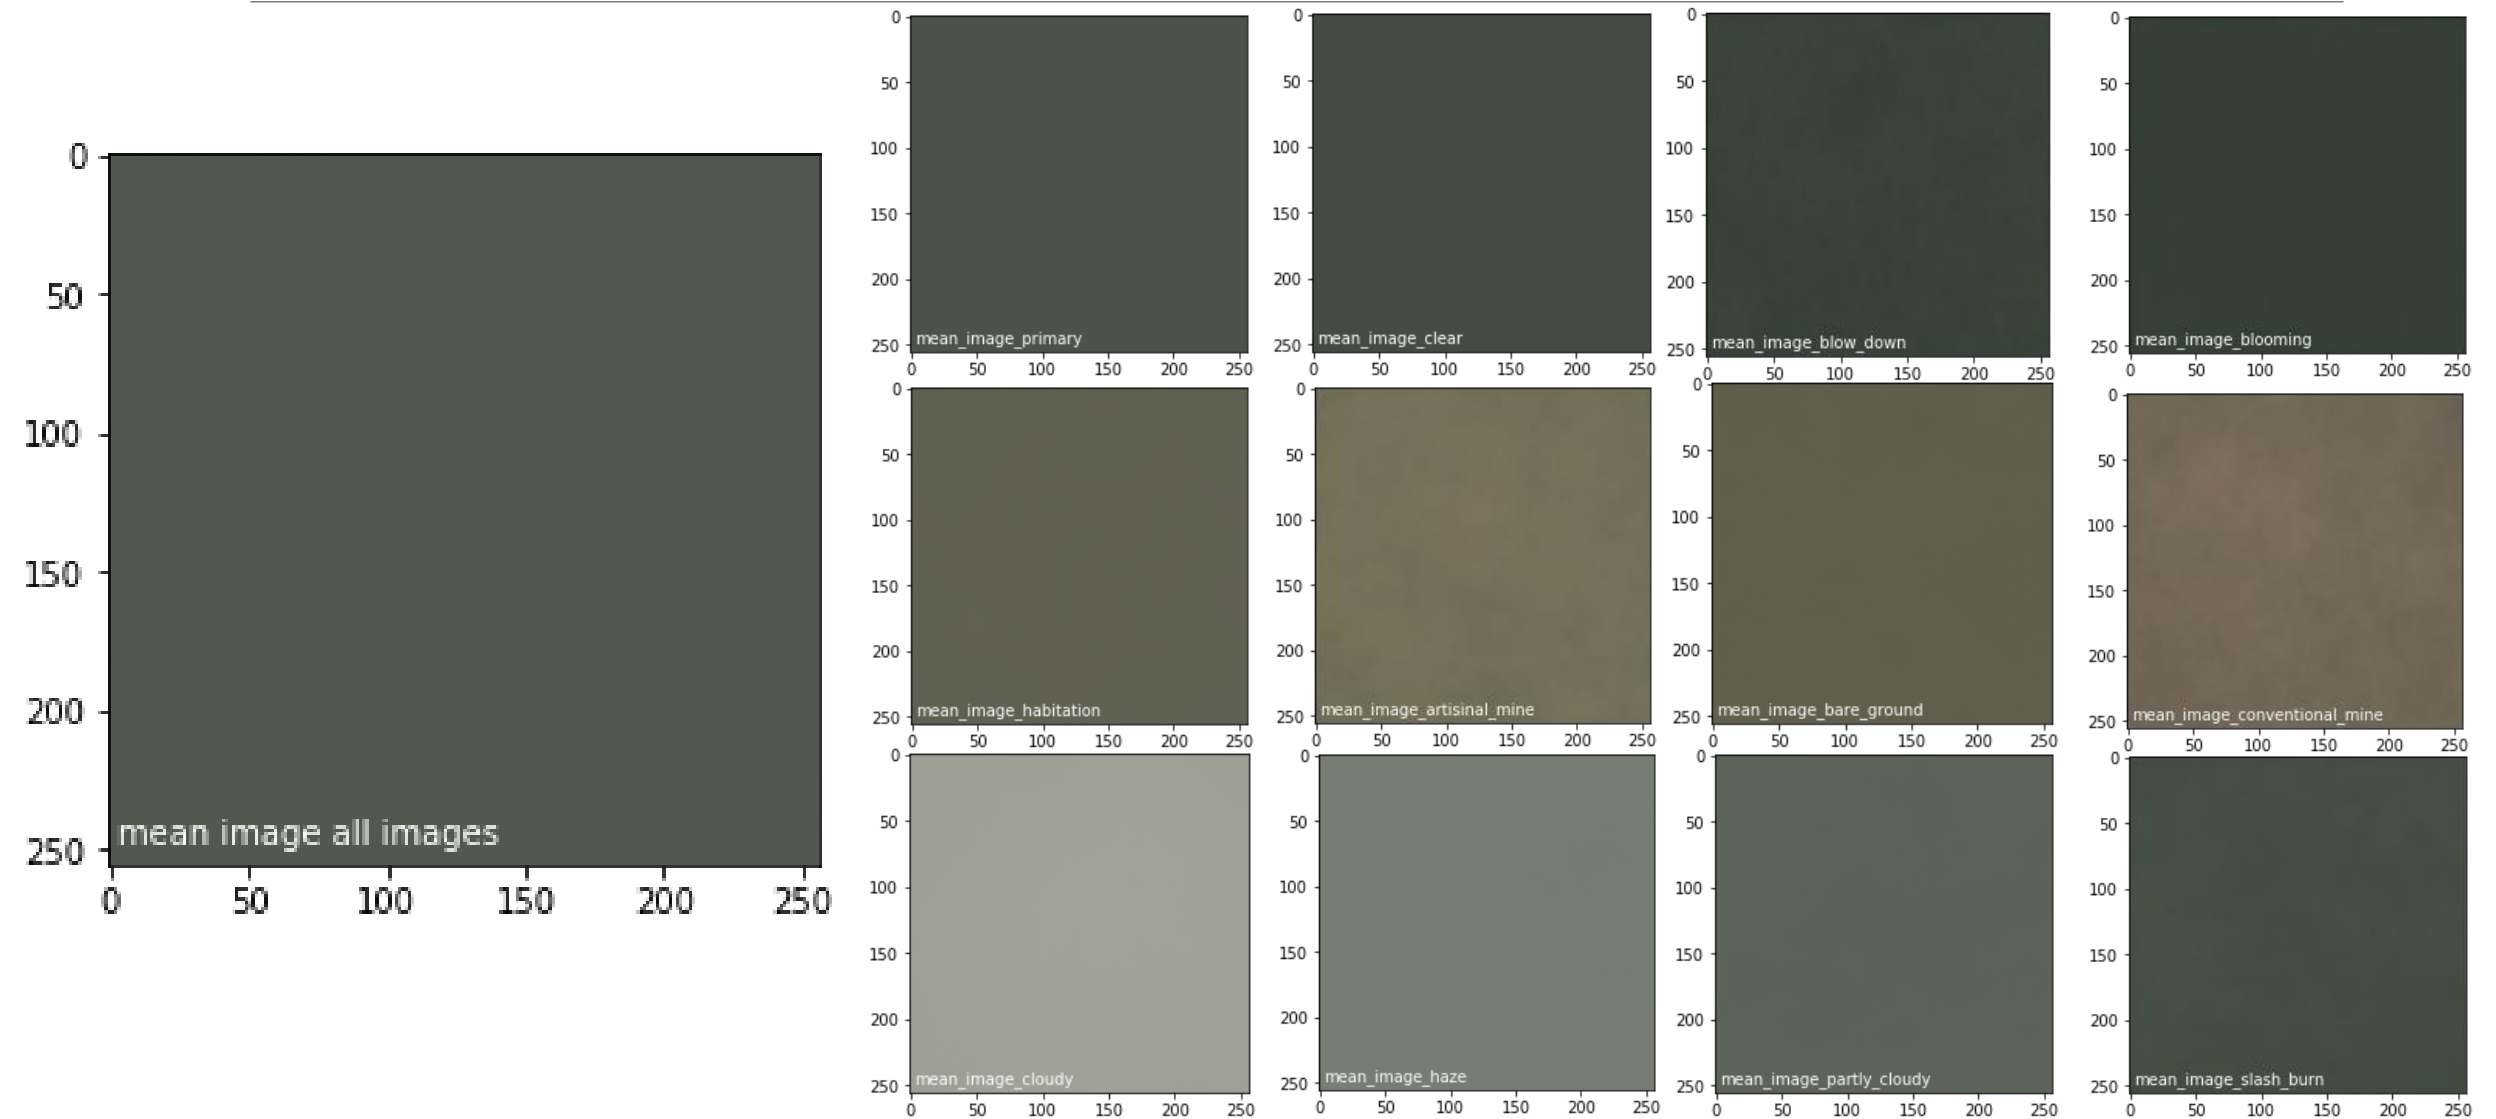
Exploratory Data Analysis – Dominant Image Colours




Exploratory Data Analysis – Dominant Image Colours (2)



Exploratory Data Analysis – Image Set Means




Exploratory Data Analysis – Image Clustering



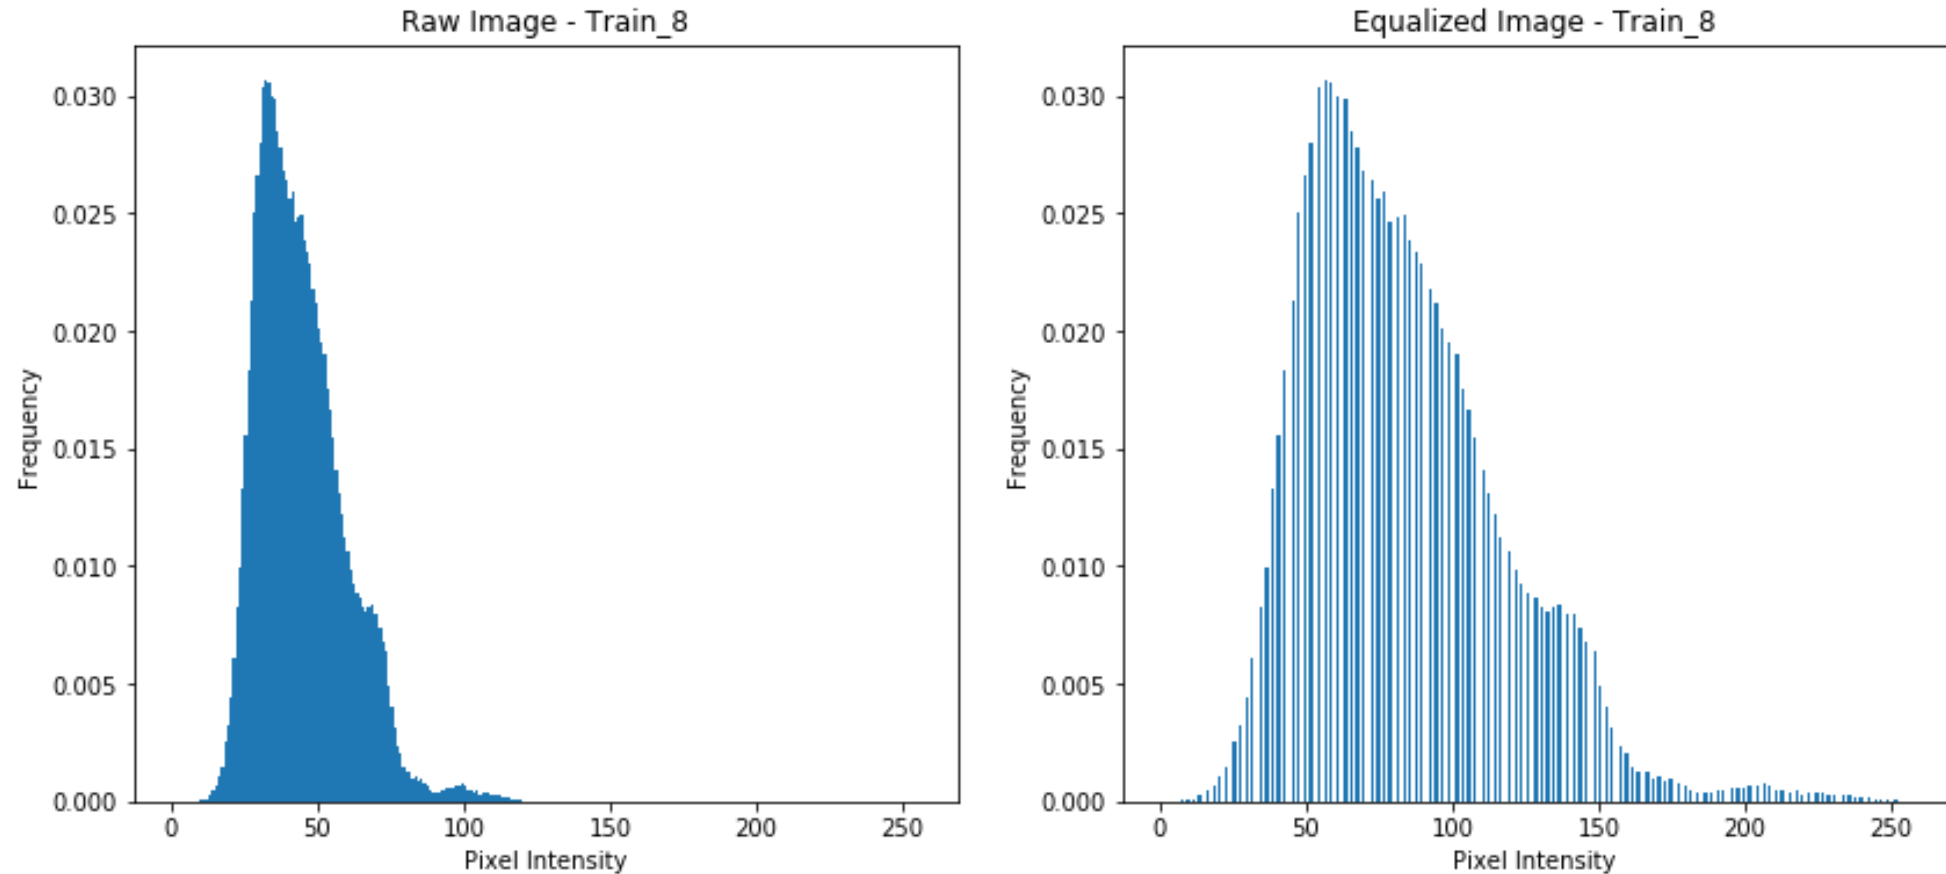
Exploratory Data Analysis – Image Clustering (2)





Exploratory Data Analysis – Image Clustering (3)



Exploratory Data Analysis – Histogram Equalization



Exploratory Data Analysis – Histogram Equalization (2)



Machine Learning – Overview

Purpose: Classify satellite images of the Amazon Rainforest as having one or more weather/land use tag.

- Convolutional Neural Networks (CNNs) are widely used to classify images.
- Several CNN models have been developed, trained, evaluated, and compared.
- Models have varying architectures.
- Implemented in Python using Keras and Tensorflow backend.

Machine Learning – Overview (2)

Purpose: Classify satellite images of the Amazon Rainforest as having one or more weather/land use tag.

Model #	Model Name	Optimizer	Regularization	Image Augmentation
1	base	SGD		
2	dropout	SGD	✓	
3	augmentation	SGD		✓
4	dropout_augment	SGD	✓	✓
5	adam	Adam		
6	adam_dropout	Adam	✓	
7	adam_augment	Adam		✓
8	adam_dropout_augment	Adam	✓	✓
9	transfer	SGD		

Machine Learning – Dataset Formatting

Dataset formatting is required to organize the raw images to be input into the CNN models.

- 1) Use dictionary of tags (keys) and integers (values) as map to create the one hot encoding sequence for each observation.
- 2) Load each image using `keras.preprocessing.image.load_img()`.
- 3) Convert each image to numpy array.
- 4) Append the one hot encoding sequence and the image numpy array to the lists of targets and images.
- 5) Array of images becomes X (train and test) and one hot encoding sequence list becomes y (train and test).
- 6) Dataset is saved in compressed format for later loading.
- 7) Once dataset is loaded, it is split into 70% train and 30% test.

Machine Learning – Model Evaluation

$$F1 = 2 * \frac{precision * recall}{(precision + recall)}$$

F1 Score: Harmonic mean of precision and recall. Equal weight for precision and recall.

Machine Learning – Model Evaluation

$$F1 = 2 * \frac{precision * recall}{(precision + recall)}$$



Introduce Beta term

$$FBeta = (1 + Beta^2) * \frac{precision * recall}{(Beta^2 * precision + recall)}$$

F1 Score: Harmonic mean of precision and recall. Equal weight for precision and recall.

Fbeta Score: Beta < 1, more weight on precision.
Beta > 1, more weight on recall.

Machine Learning – Model Evaluation

$$F1 = 2 * \frac{precision * recall}{(precision + recall)}$$



Introduce Beta term

$$FBeta = (1 + Beta^2) * \frac{precision * recall}{(Beta^2 * precision + recall)}$$



Set Beta=2

$$FBeta (F2) = 5 * \frac{precision * recall}{(4 * precision + recall)}$$

Note: Precision, recall, Fbeta calculated for each observation and then averaged. Tags with fewer observations will have less affect on the Fbeta score.

F1 Score: Harmonic mean of precision and recall. Equal weight for precision and recall.

Fbeta Score: Beta < 1, more weight on precision. Beta > 1, more weight on recall.

Multi-label precision and recall:

$$precision = \frac{1}{n} \sum_{i=1}^n \frac{|Y_i \cap h(x_i)|}{|h(x_i)|} \quad recall = \frac{1}{n} \sum_{i=1}^n \frac{|Y_i \cap h(x_i)|}{|Y_i|}$$

where:

n = # of observations

Y_i = true label assignments of the i^{th} observation




x_i = i^{th} observation

$h(x_i)$ = predicted label assignments of the i^{th} observation




Machine Learning – CNN Model Theory and Architecture

- CNNs work by extracting features from images.
- Each additional model layer increases the complexity of the learned features.
- Base model architecture:
 - 1) Images input as shape (128,128,3)
 - 2) 2x convolutional layers
 - 3) Max pooling layer
 - 4) 2x convolutional layers
 - 5) Max pooling layer
 - 6) 2x convolutional layers
 - 7) Max pooling layer
 - 8) Flattening layer
 - 9) Fully connected layer
 - 10) Output layer for prediction




Machine Learning – CNN Model Theory and Architecture

- CNNs work by extracting features from images.
- Each additional model layer increases the complexity of the learned features.
- Base model architecture:
 - 1) Images input as shape (128,128,3)
 - 2) 2x convolutional layers  **(32)** 3x3 filters, ReLU activation, 'He' weight initialization.
 - 3) Max pooling layer
 - 4) 2x convolutional layers  **(64)** 3x3 filters, ReLU activation, 'He' weight initialization.
 - 5) Max pooling layer
 - 6) 2x convolutional layers  **(128)** 3x3 filters, ReLU activation, 'He' weight initialization.
 - 7) Max pooling layer
 - 8) Flattening layer
 - 9) Fully connected layer
 - 10) Output layer for prediction

Machine Learning – CNN Model Theory and Architecture

- CNNs work by extracting features from images.
- Each additional model layer increases the complexity of the learned features.
- Base model architecture:
 - 1) Images input as shape (128,128,3)
 - 2) 2x convolutional layers
 - 3) Max pooling layer  **2x2**
 - 4) 2x convolutional layers
 - 5) Max pooling layer  **2x2**
 - 6) 2x convolutional layers
 - 7) Max pooling layer  **2x2**
 - 8) Flattening layer
 - 9) Fully connected layer
 - 10) Output layer for prediction

Machine Learning – CNN Model Theory and Architecture

- CNNs work by extracting features from images.
- Each additional model layer increases the complexity of the learned features.
- Base model architecture:
 - 1) Images input as shape (128,128,3)
 - 2) 2x convolutional layers
 - 3) Max pooling layer
 - 4) 2x convolutional layers
 - 5) Max pooling layer
 - 6) 2x convolutional layers
 - 7) Max pooling layer
 - 8) Flattening layer  **Converting to 1D vector.**
 - 9) Fully connected layer  **Shape=128, ReLU activation, 'He' initialization.**
 - 10) Output layer for prediction  **Shape=17 for 17 possible tags.**

Machine Learning – Fitting and Evaluating

- 1) Create data generator using ImageDataGenerator. Images are rescaled by a factor of $1/255$.
- 2) Use the data generator to create train and test iterators. Batch size = 128.
- 3) Fit the model using fit_generator for a set number of epochs.
- 4) Evaluate model using evaluate_generator.

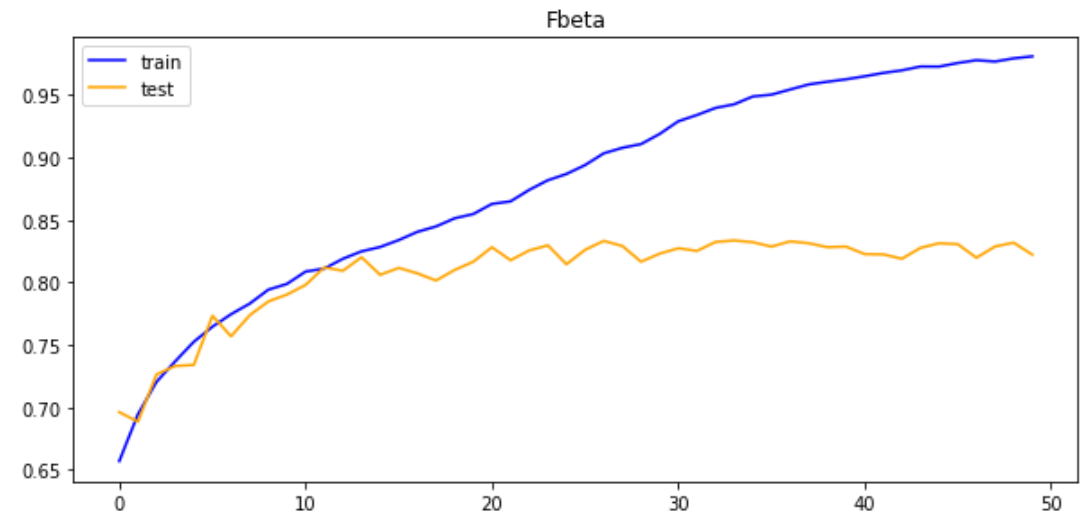

Machine Learning – Base Model

- Model architecture as described previously.

Epochs: 50

Test loss: 0.534

Test Fbeta: 0.822





Machine Learning – Dropout Regularization

- Reduce overfitting.
- Dropout layers added after each convolutional block (20% inputs dropped) and after fully connected layer (50% inputs dropped).

Epochs: 200

Test loss: 0.235

Test Fbeta: 0.855





Machine Learning – Image Data Augmentation

- Artificially increase size of training dataset by creating modified versions of raw images.
- Train data generator set to flip, rotate raw images.

Epochs: 200

Test loss: 0.126

Test Fbeta: 0.878





Machine Learning – Adam Optimizer

- Adaptive learning rate.
- Well suited for problems large in data and with noisy/sparse parameters.

Epochs: 50

Test loss: 0.711

Test Fbeta: 0.845





Machine Learning – Transfer Learning

- VGG-16 pre-trained model.
- Modified for Amazon Rainforest dataset input and predictions.

Epochs: 25

Test loss: 0.325




Test Fbeta: 0.879





Machine Learning – Model Comparison

Model #	Model Name	Optimizer	Regulariz- ation	Image Augmenta- tion	Loss	Fbeta
1	base	SGD			0.596	0.822
2	dropout	SGD	✓		0.330	0.855
3	augmentation	SGD		✓	0.138	0.878
4	dropout_augment	SGD	✓	✓	0.121	0.840
5	adam	Adam			0.910	0.845
6	adam_dropout	Adam	✓		0.345	0.862
7	adam_augment	Adam		✓	0.104	0.878
8	adam_dropout_augment	Adam	✓	✓	0.125	0.870
9	transfer	SGD			0.205	0.879


Machine Learning – Less Common Tag Prediction (1)



Machine Learning – Less Common Tag Prediction (2)



Custom Threshold vs Threshold=0.5 for Prediction:



Machine Learning – slash_burn model (1)

- **Original dataset:**
 - 209 images tagged with slash_burn.
 - 40,270 images not tagged with slash_burn.
- **Slash_burn model dataset:**
 - Train:
 - 584 raw images tagged with slash_burn:
 - 146 raw images
 - 146 raw images rotated 90°
 - 146 raw images rotated 180°
 - 146 raw images rotated 270°
 - 584 random raw images from original dataset not tagged with slash_burn.
 - Test:
 - 63 raw images tagged with slash_burn
 - 437 random raw images from original dataset not tagged with slash_burn.

Machine Learning – slash_burn model (2)

- **Original augment_model slash_burn predictive performance:**

Metric	Prediction Threshold = 0.5	Prediction Threshold = Optimized
Test Precision	0.000	0.050
Test Recall	0.000	0.506
Test Fbeta	0.000	0.182

- **Specialized slash_burn model:**

Metric	Prediction Threshold = 0.5	Prediction Threshold = Optimized
Test Precision	0.314	0.313
Test Recall	0.793	0.809
Test Fbeta	0.608	0.614

‘Guessing’ all ones achieves Test Fbeta = 0.419

Conclusions (1)

- Dropout regularization, image augmentation, and the Adam optimizer all increased model performance individually.
- The VGG-16 model was successfully used as a transfer learning model to match the performance of the best performing model built from scratch.
- The Fbeta score of all models is misleading – the metric evaluates the mean Fbeta score across all observations – the unbalanced dataset results in poorly performing tags having little effect on the overall Fbeta score.
- The Fbeta scores for less common tags are poor for all models.


Conclusions (2)

- Optimizing the prediction threshold can result in increased Fbeta scores but will result in a decrease in recall or precision. The costs and benefits should be weighed to determine if the optimized thresholds should be used.
- Using a specialized, one-vs-all model for predicting slash_burn tag, Slash_burn Fbeta was increased from 0 to 0.608 using a predictive threshold of 0.5, and from 0.182 to 0.614 when using an optimized predictive threshold.
- Image equalization did not increase the predictive performance of the specialized slash_burn model.

Further Work

- One vs all models for less common, more important tags.
- Additional data augmentation methods for expanding the datasets for less common tags.
- Test-Time augmentation.
- Ensemble model experimentation.
- Tuning learning rate.

Amazon Rainforest Image Classification



CONNOR MCANUFF - NOVEMBER 2019